

```
#Data wrangling
```

```
import pandas as pd
df = pd.read_csv('Automobile_data.csv')
df.head()
```

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	length	width	height	cu wei
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	2
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	2
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	2
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	2
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	2

```
#Identify and handle missing values
```

```
import numpy as np
```

```
# replace "?" to NaN
```

```
df.replace("?", np.nan, inplace = True)
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height	cuwei
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	2
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	2
2	1	NaN	alfa-	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	2

```
# Evaluating Missing values
missing_data = df.isnull()
missing_data.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height	curb-weight
0	False	True	False	False	False	False	False	False	False	False	False	False	False	False
1	False	True	False	False	False	False	False	False	False	False	False	False	False	False
2	False	True	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False

```
missing_data1 = df.notnull()
missing_data1.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height	curb-weight	engine-size
0	True	False	True	True	True	True	True	True	True	True	True	True	True	True	True
1	True	False	True	True	True	True	True	True	True	True	True	True	True	True	True
2	True	False	True	True	True	True	True	True	True	True	True	True	True	True	True

df.dtypes

```

symboling          int64
normalized-losses  object
make               object
fuel-type          object
aspiration         object
num-of-doors       object
body-style         object
drive-wheels       object
engine-location    object
wheel-base        float64
length            float64
width             float64
height            float64
curb-weight        int64
engine-type        object
num-of-cylinders   object
engine-size        int64
fuel-system        object
bore              object
stroke            object
compression-ratio  float64
horsepower         object
peak-rpm           object
city-mpg           int64
highway-mpg        int64
price             object
dtype: object

```

```
#Counting missing values in each column
for column in missing_data.columns.values.tolist():
    print(column)
```

```
symboling
normalized-losses
make
fuel-type
aspiration
num-of-doors
body-style
drive-wheels
engine-location
wheel-base
length
width
height
curb-weight
engine-type
num-of-cylinders
engine-size
fuel-system
bore
stroke
compression-ratio
horsepower
peak-rpm
city-mpg
highway-mpg
price
```

```
#Counting missing values in each column
for column in missing_data.columns.values.tolist():
    print(column)
    print (missing_data[column].value_counts())
    print(" ")
```

```
symboling
False      205
```

```
Name: symboling, dtype: int64
```

```
normalized-losses
```

```
False    164
```

```
True      41
```

```
Name: normalized-losses, dtype: int64
```

```
make
```

```
False    205
```

```
Name: make, dtype: int64
```

```
fuel-type
```

```
False    205
```

```
Name: fuel-type, dtype: int64
```

```
aspiration
```

```
False    205
```

```
Name: aspiration, dtype: int64
```

```
num-of-doors
```

```
False    203
```

```
True       2
```

```
Name: num-of-doors, dtype: int64
```

```
body-style
```

```
False    205
```

```
Name: body-style, dtype: int64
```

```
drive-wheels
```

```
False    205
```

```
Name: drive-wheels, dtype: int64
```

```
engine-location
```

```
False    205
```

```
Name: engine-location, dtype: int64
```

```
wheel-base
```

```
False    205
```

```
Name: wheel-base, dtype: int64
```

```
length
```

```
False      205
Name: length, dtype: int64
```

```
width
False      205
Name: width, dtype: int64
```

```
height
False      205
Name: height, dtype: int64
```

```
curb-weight
False      205
Name: curb-weight, dtype: int64
```

```
#"normalized-losses","stroke","bore","horsepower","peak-rpm"  replace by mean or median (numeric data)
avg_1 = df["normalized-losses"].astype("float").mean()
avg_1
```

```
122.0
```

```
df["normalized-losses"].replace(np.nan, avg_1, inplace = True)
```

```
avg_2=df['bore'].astype('float').mean()
avg_2
```

```
3.3297512437810957
```

```
df['bore'].replace(np.nan, avg_2, inplace= True)
```

```
avg_3 = df['stroke'].astype('float').mean(axis=0)
df['stroke'].replace(np.nan, avg_3, inplace = True)
```

```
avg_4=df['horsepower'].astype('float').mean(axis=0)
```

```
df['horsepower'].replace(np.nan, avg_4, inplace= True)
```

```
avg_5=df['peak-rpm'].astype('float').mean(axis=0)  
df['peak-rpm'].replace(np.nan, avg_5, inplace= True)
```

```
#replace by mode or maximum occurring frequency  
df['num-of-doors'].value_counts()
```

```
four      114  
two        89  
Name: num-of-doors, dtype: int64
```

```
df['num-of-doors'].value_counts().idxmax()
```

```
'four'
```

```
#replace the missing 'num-of-doors' values by the most frequent  
df["num-of-doors"].replace(np.nan, "four", inplace = True)
```

```
# simply drop whole row with NaN in "price" column  
df.dropna(subset=["price"], axis=0, inplace = True)
```

```
df.reset_index(drop = True, inplace = True)
```

```
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	width	height	cu wei
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	2
1	3	122	alfa-	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	2

df.dtypes

```

symboling          int64
normalized-losses  object
make              object
fuel-type         object
aspiration        object
num-of-doors      object
body-style        object
drive-wheels      object
engine-location   object
wheel-base       float64
length            float64
width             float64
height            float64
curb-weight       int64
engine-type       object
num-of-cylinders  object
engine-size       int64
fuel-system       object
bore              object
stroke            object
compression-ratio float64
horsepower        object
peak-rpm          object
city-mpg          int64
highway-mpg       int64
price             object
dtype: object

```


Data Standardization

```
df['city-l/100km']=235/df["city-mpg"]
```

```
df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")  
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")  
df[["price"]] = df[["price"]].astype("float")  
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

```
df.dtypes
```

```
#data transformation for highway-mpg into L/100 km  
#data normalization :scaling within 1  
df['length'] = df['length']/df['length'].max()  
df['width'] = df['width']/df['width'].max()
```

```
df['height'] = df['height']/df['height'].max()  
df[["length","width","height"]].head()
```

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

```
#Binning
```

```
df["horsepower"]=df["horsepower"].astype(float)
```

```
df["horsepower"]
```

```
0      111.0
1      111.0
2      154.0
3      102.0
4      115.0
```

```
...
196     114.0
197     160.0
198     134.0
199     106.0
200     114.0
```

```
Name: horsepower, Length: 201, dtype: float64
```

```
binwidth = (max(df["horsepower"])-min(df["horsepower"]))/4
```

```
binwidth
```

```
53.5
```

```
bins = np.arange(min(df["horsepower"]), max(df["horsepower"]), binwidth)
```

```
bins
```

```
array([ 48. , 101.5, 155. , 208.5])
```

```
group_names = ['Low', 'Medium', 'High']
```


```
df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,include_lowest=True )
```

```
df[['horsepower', 'horsepower-binned']].head(20)
```

	horsepower	horsepower-binned
0	111.0	Medium
1	111.0	Medium
2	154.0	Medium
3	102.0	Medium
4	115.0	Medium
5	110.0	Medium
6	110.0	Medium
7	110.0	Medium
8	140.0	Medium
9	101.0	Low
10	101.0	Low
11	121.0	Medium
12	121.0	Medium
13	121.0	Medium
14	182.0	High
15	182.0	High
16	182.0	High
17	48.0	Low

```
from matplotlib import pyplot as plt
plt.hist(df["horsepower"], bins = 3)
plt.xlabel("horsepower")
plt.ylabel("count")
plt.title("horsepower bins")
```

```
Text(0.5, 1.0, 'horsepower bins')
```



```
#Indicator variable
```

```
df.columns
```

```
Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
      'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
      'highway-mpg', 'price', 'horsepower-binned'],
      dtype='object')
```

```
dummy_variable_1 = pd.get_dummies(df["fuel-type"])
```

```
dummy_variable_1.rename(columns={'fuel-type-diesel':'gas', 'fuel-type-diesel':'diesel'}, inplace=True)
dummy_variable_1.head()
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1

```
df = pd.concat([df, dummy_variable_1], axis=1)
df.drop("fuel-type", axis = 1, inplace=True)
```

```
df.head()
```

	symboling	normalized- losses	make	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	length	width	height	cur weig
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	0.890278	0.816054	25
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	0.890278	0.816054	25
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	0.909722	0.876254	28
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	0.919444	0.908027	23
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	0.922222	0.908027	28

```
dummy_variable_2 = pd.get_dummies(df['aspiration'])
dummy_variable_2.rename(columns={'std': 'aspiration-std', 'turbo': 'aspiration-turbo'}, inplace=True)
dummy_variable_2.head()
```

	aspiration-std	aspiration-turbo
0	1	0
1	1	0
2	1	0
3	1	0

```
df = pd.concat([df, dummy_variable_2], axis=1)
df.drop('aspiration', axis = 1, inplace=True)
```

```
df.to_csv('clean_df.csv')
```

